

LAB 1: AI-ASSISTED CITATION INTEGRITY AUDIT

Generate a Research Paper with AI, Then Audit Its Citations

AI Tools for Research • 2 Hours • Individual Activity • Topic IDs T1-T40

Student Detail	Fill in
Name	Vibhum Sharma
Roll Number	RSI2026506
Programme / Department	Ph.D. IT
Topic ID	T32
Full Assigned Topic	Orchestrating the Virtual Tumor Board: A Scoping Review of Multi-Agent AI and Collaborative Decision-Making in Clinical Oncology
AI Tool Used	Google Gemini 3.1 Pro Extended
Date	21 August 2026

IMPORTANT: You will receive a topic from T1-T40. Your task is not to produce the “best” paper. Your task is to generate an AI-written research paper and systematically determine whether its citations can be trusted.

1. PURPOSE OF THIS LAB

Generative AI systems can produce research papers that look scholarly and may contain convincing references, author names, journal titles, DOIs, arXiv IDs, and in-text citations. However, the presence of a citation does not guarantee that the cited publication actually exists, that its metadata are correct, that the identifier belongs to that publication, or that the cited paper supports the statement made by the AI.

In this laboratory activity, you will test these issues yourself by generating a research paper and auditing a systematic sample of its references.

Question 1 - Citation Authenticity: Does the cited publication actually exist, and is it correctly identified?

Question 2 - Claim-Citation Support: If the publication is genuine, does it actually support the claim for which the AI cited it?

2. IMPORTANT RULES

- Use only the topic assigned to you by the instructor.
- Generate the paper in a fresh AI conversation.
- Do not tell the AI that you are conducting a citation audit or hallucination test.
- Do not intentionally ask the AI to generate fake references.
- Preserve the original AI-generated paper exactly as produced.
- Do not correct, regenerate, replace, or question references before saving the original output.
- For verification, do not accept another AI chatbot's statement as proof that a reference exists.
- Use scholarly databases, publisher pages, DOI/arXiv/PMID records, or other reliable scholarly sources.
- Record what you actually find, even if all references are genuine or all are problematic.

Most important principle: You are not being evaluated on whether the AI produces good or bad references. You are being evaluated on how systematically and correctly you audit them.

3. WHAT YOU WILL DO

Part	Activity
A	Record your baseline perception of AI citations
B	Record the AI environment used
C	Generate and preserve an AI-written research paper
D	Create a reference inventory
E	Select 10 references systematically
F	Judge their plausibility before verification
G	Verify authenticity, metadata and identifiers
H	Determine whether genuine references support the AI-generated claims

4. RECOMMENDED 2-HOUR TIME PLAN

Time	Activity
0-10 min	Parts A-B: Baseline and AI environment
10-30 min	Part C: Generate and save research paper
30-40 min	Parts D-E: Inventory and select references
40-45 min	Part F: Pre-verification plausibility
45-80 min	Part G: Authenticity and metadata audit
80-110 min	Part H: Claim-citation entailment audit
110-120 min	Calculate scores, check worksheet and prepare submission

Do not spend excessive time formatting the AI-generated paper. The audit is the main activity.

PART A - BASELINE ASSESSMENT

Complete this section BEFORE opening the AI tool.

A1. How often do you use generative AI for academic work?

☐ Never ☐ Rarely ☐ Occasionally ☒ Weekly ☐ Almost daily

A2. Have you previously asked an AI system to generate references?

☐ Frequently ☐ Occasionally ☒ Once or twice ☐ Never

A3. Have you previously checked whether an AI-generated citation actually exists?

☐ Always ☐ Sometimes ☐ Rarely ☒ Never ☐ Not applicable

A4. Before this lab, how confident are you that citations generated by AI are generally reliable?

☐ 1 ☐ 2 ☒ 3 ☐ 4 ☐ 5 (1 = not at all confident; 5 = completely confident)

Important: Do not change this answer later, even if your opinion changes during the experiment.

PART B - RECORD YOUR AI ENVIRONMENT

Information	Your Response
AI tool used	Google Gemini
Model name	Pro
Model/version shown	3.1
Access tier	PRO
Web browsing available?	YES
Was web browsing actually used?	NOT SURE
Research/Deep Research mode used?	YES
Date/time generated	August 21, 2026 [09:40 AM]

Important: If the model/version is not displayed, write NOT SHOWN. Do not guess the model version.

PART C - GENERATE THE AI RESEARCH PAPER

C1. Open a Fresh AI Conversation

Do not tell the AI that this is a citation audit, that you are testing hallucinations, or that you expect fabricated references. We want to observe normal AI research-paper generation behaviour.

C2. Use the Following Common Prompt

Replace only [YOUR ASSIGNED TOPIC].

Generate a scholarly research paper on the topic: "[YOUR ASSIGNED TOPIC]". The paper should be suitable for an academic audience and should not exceed approximately 10 pages. Include Title, Abstract, Keywords, Introduction, Background/Related Work/Literature Review, Main Thematic or Technical Discussion, Current Approaches or Developments, Research Challenges and Limitations, Research Gaps and Future Directions, Conclusion, and Complete References. Use scholarly literature to support factual, technical, comparative and research claims. Provide in-text citations throughout the paper. For every reference, provide as much bibliographic information as available, including authors, title, publication year, journal/conference/source, and DOI, PMID, arXiv ID or another persistent identifier wherever available. Use genuine scholarly sources. Do not intentionally fabricate references.

C3. Preserve the Original Output

- ☐ Save the complete AI-generated paper
- ☐ Save the complete reference list
- ☐ Take screenshot(s) showing the generated output/references
- ☐ Save conversation/export/share evidence if available

STOP: Do not ask the AI to correct or verify anything yet. Your original output is the experimental evidence.

PART D - PAPER PROFILE AND REFERENCE INVENTORY

Paper title: Orchestrating the Virtual Tumor Board: A Scoping Review of Multi-Agent AI and Collaborative Decision-Making in Clinical Oncology

Measure	Value
Approximate pages	7
Approximate word count	1940
Number of sections	7 Numbered + 2 Abstract + Reference Section
Total number of references	7
Approximate in-text citation occurrences	14-16 but has no explicit in-text citation marker anywhere in the body text.
Did AI warn you to verify references?	NO

D1. Reference Inventory

Classify every reference once using the best available identifier in this priority order: DOI -> PMID -> arXiv ID -> URL only -> No persistent identifier. The category totals must equal the total number of references in the AI-generated paper.

Reference #	Author/Year	Identifier	Category
1	Powell et al., 2021	DOI 10.1200/op.20.00480	DOI available
2	Ekhtator et al., 2022	DOI 10.7759/cureus.25682	DOI available
3	Thiagarajan et al., 2023	DOI 10.1200/go.22.00348	DOI available
4	Park & Chae, 2023	DOI 10.1007/s12029-023-00967-8	DOI available
5	Wang, X. (n.d.)	none supplied	No persistent identifier
6	Deniz, A. Z. (n.d.)	none supplied	No persistent identifier
7	Liu, Q. (n.d.) — "EvoMDT"	none supplied	No persistent identifier

Reference Type	Count
DOI available	4
PMID available, no DOI	0
arXiv ID available, no DOI/PMID	0
URL only	0
No persistent identifier supplied	3
TOTAL	7

Examples

Category	Example	How to Count
DOI available	Kather, J. N., et al. (2019). Predicting survival from colorectal cancer histology slides using deep learning. PLoS Medicine. DOI: 10.1371/journal.pmed.1002730	DOI available
PMID available, no DOI	A biomedical paper for which the AI provides PMID: 12345678, but no DOI is supplied/found	PMID available, no DOI
arXiv ID available	Yao, S., et al. ReAct: Synergizing Reasoning and Acting in Language Models. arXiv:2210.03629	arXiv ID available
URL only	World Health Organization. Colorectal cancer. Available at: https://... with no DOI, PMID or arXiv ID	URL only
No persistent identifier	Goodfellow, I., Bengio, Y., & Courville, A. (2016). Deep Learning. MIT Press.	No persistent identifier if none is supplied

Important: Absence of a DOI does not mean a reference is fake. An identifier being present also does not prove authenticity; it must be checked.

PART E - SELECT 10 REFERENCES FOR AUDIT

You are not allowed to choose only suspicious-looking references. This prevents selection bias.

If the bibliography contains 10 or fewer references: audit all references.

If it contains more than 10 references: select the first 3 references, the last 3 references, and 4 references approximately evenly distributed through the middle.

Example: If there are 28 references, a reasonable sample could be 1, 2, 3, 8, 13, 18, 23, 26, 27, 28.

Audit Slot	Original Bibliography Reference No.
R1	1
R2	2
R3	3
R4	4
R5	5
R6	6
R7	7

Important: From this point onward, R1-R10 refer to your audit sample, not necessarily bibliography references 1-10.

PART F - PLAUSIBILITY TEST

Do this BEFORE searching for any reference.

Look only at the citation as generated by the AI. Ask yourself: “If I had not verified this citation, how convincing would it look to me?”

Score	Meaning
-------	---------

1	Obviously suspicious
2	Somewhat suspicious
3	Unsure
4	Looks genuine
5	Completely convincing

Ref.	Initial Impression	Plausibility 1-5	Predict Genuine?
R1	looks like a real article, authors are there, doi is there, journal name is there	5	YES
R2	looks like a real cureus article, authors are there, doi is there	5	YES
R3	looks like a real article, authors are there, doi is there, journal name is there	5	YES
R4	looks like a real article, authors are there, doi is there, journal name is there	5	YES
R5	Only title, no author name or journal, not even DOI	2	NO
R6	Only title, no author name or journal, not even DOI	2	NO
R7	Only title, no author name or journal, not even DOI	3	NO

Example: A citation with realistic authors, a well-known journal, complete volume/pages and a DOI-like identifier may look highly convincing. You might record Plausibility = 5 and Predict Genuine = Yes. You are recording your pre-verification belief, not the truth.

Critical rule: Never change Part F after completing Part G. Your incorrect predictions are valuable experimental observations.

PART G - CITATION AUTHENTICITY AND METADATA AUDIT

G1. What Are You Checking?

- Does the publication exist?
- Is the title correct?
- Are the authors correct?
- Is the year correct?
- Is the journal/conference/source correct?
- Are volume, issue and pages correct, if supplied?
- Does the DOI/arXiv/PMID belong to this publication?

G2. Verification Procedure

- ☐ Persistent identifier: DOI / arXiv / PMID
- ☐ Publisher or official repository
- ☐ Google Scholar
- ☐ Semantic Scholar / OpenAlex
- ☐ Crossref
- ☐ PubMed or another discipline-specific database
- ☐ Exact-title search

☐ Author + title-keyword search

Not acceptable as sole evidence: “I asked another AI and it said the paper exists.” AI can help you search, but final verification must rely on scholarly or publisher evidence.

G3. Assign One A-E Code

Code	Classification	Meaning
A	VERIFIED	Publication exists and important metadata are correct.
B	WRONG METADATA	Publication exists, but one or more bibliographic details are incorrect.
C	FRANKENSTEIN	Real-looking pieces are combined, but that specific publication does not exist.
D	FABRICATED	No reliable scholarly evidence that the publication exists.
E	IDENTIFIER MISMATCH	DOI/arXiv/PMID is invalid or resolves to a different publication.

G4. Examples

- A:** AI gives a paper and all details match the publisher record.
- B:** The paper exists, but the AI gives the wrong publication year or journal.
- C:** The authors are real, the journal is real, and the topic is plausible, but no such paper with that title/authorship combination exists.
- D:** You cannot find the paper or a credible scholarly record after systematic searching.
- E:** The DOI supplied by AI opens an entirely different paper or does not resolve.

G5. Record Your Evidence

Ref.	Publication Found?	Title Match?	Authors Match?	Year Match?	Venue Match?	Identifier Match?
R1	Y	Y	Y	Y	Y	Y
R2	Y	Y	Y	Y	Y	Y
R3	Y	Y	Y	Y	Y	Y
R4	Y	Y	Y	Y (epub year; print issue is 2024)	Y	Y
R5	Y	Y	N (1 of 6 authors given)	N (none given; actual 2025)	N (none given; actual iScience)	N/A (none supplied)
R6	Y	Y	N (1 of 8 authors given)	N (none given; actual 2026)	N (none given; actual Updates in Surgery)	N/A (none supplied)
R7	Y	Y	N (1 of 13 authors given)	N (none given; actual 2026)	N (none given; actual npj Digital Medicine)	N/A (none supplied)

Ref.	Verification Source(s)	Identifier Checked	What You Found
R1	PubMed, ASCO Publications (ascopubs.org), Crossref DOI resolver	DOI: 10.1200/OP.20.00480	Fully matches — PMID 33031011, <i>JCO Oncol Pract</i> 17(6), e809–e816, 2021
R2	PubMed, PMC, Cureus journal site	DOI: 10.7759/cureus.25682	Fully matches — PMID 35677741, <i>Cureus</i> 14(6):e25682, 2022
R3	PubMed, ASCO Publications, PMC	DOI: 10.1200/GO.22.00348	Fully matches — PMID 36706349, <i>JCO Glob Oncol</i> 9:e2200348, 2023
R4	PubMed, PMC, Springer	DOI: 10.1007/s12029-023-00967-8	Fully matches — PMID 37702851, <i>J Gastrointest Cancer</i> 55(1):365–372. Epub Sep 2023, print issue Mar 2024 — AI used epub year; immaterial discrepancy
R5	Springer (as cited elsewhere), Nature/ScienceDirect (<i>iScience</i>), PMC	DOI: 10.1016/j.isci.2025.114082 (<i>found; none supplied by AI</i>)	Real paper exists, titled and topically identical, but AI dropped 5 of 6 authors (Wang Q, Ding G, Wang J, Tang Y, Feng Y) and gave no year, journal, or DOI
R6	PubMed, Springer (<i>Updates in Surgery</i>)	DOI: 10.1007/s13304-026-02749-w (<i>found; none supplied by AI</i>)	Real paper exists (<i>Updates Surg</i> , 2026, PMID 42377710), but AI dropped 7 of 8 authors and gave no year, journal, or DOI
R7	Nature (<i>npj Digital Medicine</i>), ResearchGate, dblp	DOI: 10.1038/s41746-025-02304-8 (<i>found; none supplied by AI</i>)	Real paper exists (<i>npj Digital Medicine</i> 9:124, 2026), but AI dropped 12 of 13 authors and gave no year, journal, or DOI

G6. Final Classification

Ref.	Final Code	One-Line Reason
R1	A	All bibliographic details (authors, journal, year, DOI, PMID) match the ASCO/PubMed record exactly
R2	A	All details match Cureus/PubMed record exactly, including PMID
R3	A	All details match ASCO/PubMed record exactly, including PMID

R4	A	All core details match; only trivial epub-vs-print year variance, not a factual error
R5	B	Real <i>iScience</i> (2025) paper, but AI listed only 1 of 6 authors and omitted year, journal, and DOI entirely
R6	B	Real <i>Updates in Surgery</i> (2026) paper, but AI listed only 1 of 8 authors and omitted year, journal, and DOI entirely
R7	B	Real <i>npj Digital Medicine</i> (2026) EvoMDT paper, but AI listed only 1 of 13 authors and omitted year, journal, and DOI entirely

Good reason: B: Article exists, but AI reported 2022 instead of 2020.

Poor reason: B: Details wrong.

PART G - SCORING

Count your final classifications: A = 4 B = 3 C = 0 D = 0 E = 0

Total audited: N = A + B + C + D + E = 7

Code	Points
A	4
B	3
C	1
D	0
E	1

Authenticity Points = $4A + 3B + C + E$

Authenticity Score = $[(4A + 3B + C + E) / (4N)] \times 100$

My calculation: $4(4) + 3(3) + 1(0) + 0(0) + 1(0) = 25$ points

Authenticity Score = $25 / (4 \times 7) \times 100 = 89.3/100$

Worked example: If A=6, B=2, C=1, D=1, E=0, then points = $4(6)+3(2)+1(1)=31$; maximum = 40; Authenticity Score = $31/40 \times 100 = 77.5$.

Score	Interpretation
90-100	Excellent authenticity
75-89	Good, minor concerns
60-74	Moderate citation-authenticity risk
40-59	High citation-authenticity risk
Below 40	Very high citation-authenticity risk

G7. Prediction Accuracy

Compare your Part F Genuine Yes/No predictions with the Part G results.

Ref	Predicted	Actually Genuine?	Correct
R1	Yes	Yes	✓
R2	Yes	Yes	✓
R3	Yes	Yes	✓
R4	Yes	Yes	✓
R5	No	Yes	✗
R6	No	Yes	✗

R7	No	Yes	✗
----	----	-----	---

Prediction Accuracy = (Correct genuine/fake predictions / Total predictions) x 100

Correct predictions = 4 Total predictions = 7 Prediction Accuracy = 57.1 %

FINAL LAB 1 RESULTS

Measure	Your Result
Total references in AI paper	7
References deeply audited	7
A: Verified	4
B: Wrong metadata	3
C: Frankenstein	0
D: Fabricated	0
E: Identifier mismatch	0
Authenticity Score	89.3/100
Prediction Accuracy	57.1 %

FINAL INTERPRETATION

1. Did professional-looking citations necessarily turn out to be genuine?

Not entirely, but not in the way I expected going in. All 7 references turned out to correspond to real, existing publications — nothing was fabricated. But 3 of the 7 (R5, R6, R7) *looked* the least professional in Part F — missing authors, year, journal, and DOI — and I predicted them as fake. They weren't. The AI had simply stripped almost all identifying metadata while keeping the title exact. So "looks professional" wasn't a reliable signal of genuineness in either direction: the 4 fully-formed citations were fully genuine, but the 3 stripped-down ones were also genuine, just badly reported. Professional appearance predicted metadata completeness, not existence.

2. Did all genuine references actually support the claims for which they were cited?

I didn't run a formal Part H claim-by-claim check since it isn't a separate section in this worksheet, but based on reading the paper against what I verified: yes, broadly. R1–R4 support the specific numeric/clinical claims they're attached to (e.g., the VTB-adoption and concordance-rate figures line up with each paper's actual reported results). R5–R7 are cited in more general, summary-style sentences ("comprehensive reviews... indicate that AI tools significantly enhance...") rather than specific quantitative claims, so their support is harder to pin down precisely — they're topically relevant but the paper doesn't attribute a specific number or finding to any one of them individually.

3. What was the most serious citation-integrity problem you observed?

Not fabrication — the more serious problem was **silent metadata stripping on real papers**. For R5–R7, the AI had clearly "seen" the correct title (word-for-word matches to real papers, including a 2026 paper published after training-relevant dates, suggesting live retrieval), but discarded 83–92% of the actual author list and every other bibliographic field, without flagging that it was doing so. A reader has no way to tell this citation apart from a partially-remembered or fabricated one just by looking at it — which is exactly the trap Part F was designed to expose. That's arguably worse than obvious fabrication, because it's harder to catch and easier to trust.

4. What is the single most important lesson you learned from this lab?

Citation risk from AI isn't binary (real vs. fake) — it's a spectrum, and the most dangerous point on that spectrum is a citation that's real but incompletely or lazily reported, because it passes a naive "does this paper exist?" check while still being unusable as an actual scholarly reference (you can't locate, verify, or cite it properly from what's given). My pre-verification instinct treated "missing metadata" as a

fabrication signal, and it cost me 3 out of 7 predictions. Going forward, the practical takeaway is: incomplete citation $\neq$ fake citation, but it does mean don't trust the citation until you've independently pulled the full record yourself.

SUBMISSION INSTRUCTIONS

Submit one PDF: Lab1_RollNumber_Name

Example: Lab1_IEC2026001_RahulSharma

STUDENT DECLARATION

I confirm that I generated the research paper using the recorded AI environment, preserved the original output before conducting the audit, followed the prescribed reference-selection procedure, and personally verified the evidence reported in this worksheet. I have not altered my pre-verification predictions after seeing the verification results.

Student Name: VIBHUM SHARMA

Roll Number: RSI2026506

Signature: Vibhum Sharma

Date: August 21, 2026

CENTRAL LESSON OF LAB 1

Does the source exist? + Does the source actually support the claim?

AI Generation -> Reference Inventory -> Systematic Sampling -> Prediction -> Authenticity Verification -> Claim-Support Verification -> Quantitative Scores -> Submission